

## Classification with BHI Dataset and VGG-style network

In this experiment you will set up a VGG-style network to classify histopathologic scans of breast tissue from the [BHI](https://www.kaggle.com/paultimothymooney/breast-histopathology-images) (<https://www.kaggle.com/paultimothymooney/breast-histopathology-images>) dataset.

```
In [1]: ▶ import tensorflow.keras as keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten, MaxPooling2D
from tensorflow.keras.optimizers import SGD, Adam
from matplotlib import pyplot as plt
import numpy as np
```

WARNING:tensorflow:From C:\Users\Johan\anaconda3\envs\py311\Lib\site-packages\keras\src\losses.py:2976: The name tf.losses.sparse\_softmax\_cross\_entropy is deprecated. Please use tf.compat.v1.losses.sparse\_softmax\_cross\_entropy instead.

Here we use a Keras utility function to load the dataset. I already organized the data into HDF5 files which are a good format for storing array data.

```
In [2]: ▶ from tensorflow.keras.utils import get_file
x_train_path = get_file('idc_train.h5', 'https://storage.googleapis.com/data401-datasets/idc_train.h5')
x_test_path = get_file('idc_test.h5', 'https://storage.googleapis.com/data401-datasets/idc_test.h5')
```

We read the data from the HDF5 files into Numpy arrays.

I crop the images so they are all 48x48.


```
In [3]: ▶ import h5py as h5
with h5.File(x_train_path, 'r') as f:
    x_train = f['X'][:, :, 1:49, 1:49] # Load half the data to avoid out-of-memory errors
    y_train = f['y'][:, :2]
with h5.File(x_test_path, 'r') as f:
    x_test = f['X'][:, :, 1:49, 1:49]
    y_test = f['y'][:, :]
```

```
In [4]: ▶ x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
Out[4]: ((61925, 48, 48, 3), (61925,), (13761, 48, 48, 3), (13761,))
```

Showing a few images from the dataset.

```
In [5]: ▶ for i in range(5):
plt.imshow(np.squeeze(x_train[i]))
plt.title(y_train[i])
plt.show()
```



## Data preprocessing

1. Convert the train and test images to floating point and divide by 255.
2. Compute the average value of the entire training image set.
3. Subtract the average value from the training and testing images.

```
In [6]: ► x_test = x_test.astype("float32")/255  
x_train = x_train.astype("float32")/255  
  
x_train_avg = np.mean(x_train)  
  
x_train = x_train-x_train_avg  
x_test = x_test-x_train_avg
```

Build a VGG-style binary classifier model. For example, your network could contain the following:

1. 32 convolutional filters of size 3x3, zero padding, ReLU activation
2. 2x2 max pooling with stride 2
3. 64 filters
4. max pool
5. 128 filters
6. max pool
7. 256 filters
8. max pool
9. flatten
10. Fully-connected layer with 128 outputs
11. Final binary classification layer

```
In [7]: ▶ model = Sequential([
    Input(x_train.shape[1:]),
    Conv2D(32,3,activation='relu',padding='same',name='conv1'),
    MaxPooling2D(2,2),
    Conv2D(64,3,activation='relu',padding='same',name='conv2'),
    MaxPooling2D(2,2),
    Conv2D(128,3,activation='relu',padding='same',name='conv3'),
    MaxPooling2D(2,2),
    Conv2D(256,3,activation='relu',padding='same',name='conv4'),
    MaxPooling2D(2,2),
    Flatten(),
    Dense(128,activation='relu',name='dense1'),
    Dense(2,activation='softmax',name='z')
])
model.summary()
```

WARNING:tensorflow:From C:\Users\Johan\anaconda3\envs\py311\Lib\site-packages\keras\src\backend.py:1398: The name tf.executing\_eagerly\_outside\_functions is deprecated. Please use tf.compat.v1.executing\_eagerly\_outside\_functions instead.

WARNING:tensorflow:From C:\Users\Johan\anaconda3\envs\py311\Lib\site-packages\keras\src\layers\pooling\max\_pooling2d.py:161: The name tf.nn.max\_pool is deprecated. Please use tf.nn.max\_pool2d instead.

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv1 (Conv2D)	(None, 48, 48, 32)	896
max_pooling2d (MaxPooling2D)	(None, 24, 24, 32)	0
conv2 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
conv3 (Conv2D)	(None, 12, 12, 128)	73856
max_pooling2d_2 (MaxPooling2D)	(None, 6, 6, 128)	0
conv4 (Conv2D)	(None, 6, 6, 256)	295168
max_pooling2d_3 (MaxPooling2D)	(None, 3, 3, 256)	0
flatten (Flatten)	(None, 2304)	0
dense1 (Dense)	(None, 128)	295040
z (Dense)	(None, 2)	258

=====

Total params: 683714 (2.61 MB)  
 Trainable params: 683714 (2.61 MB)

Non-trainable params: 0 (0.00 Byte)

---

Set up the model to optimize the sparse categorical cross-entropy loss using Adam optimizer and learning rate of .0003. Calculate accuracy metrics during training.

In [8]: ▶

```
learning_rate = 3e-4

opt = Adam(learning_rate=learning_rate)

model.compile(loss='sparse_categorical_crossentropy', optimizer=opt, metrics='accuracy')
```

Now fit the model to the data using a batch size of 32 and 10% validation split over 10 epochs.

```
In [9]: ▶ batch_size = 32  
epochs = 10  
  
history = model.fit(x_train,y_train,batch_size=batch_size,epochs=epochs,validation_split=0.1,verbose=True)
```

Epoch 1/10

WARNING:tensorflow:From C:\Users\Johan\anaconda3\envs\py311\Lib\site-packages\keras\src\utils\tf\_utils.py:492: The name tf.ragged.RaggedTensorValue is deprecated. Please use tf.compat.v1.ragged.RaggedTensorValue instead.

WARNING:tensorflow:From C:\Users\Johan\anaconda3\envs\py311\Lib\site-packages\keras\src\engine\base\_layer\_utils.py:384: The name tf.executing\_eagerly\_outside\_functions is deprecated. Please use tf.compat.v1.executing\_eagerly\_outside\_functions instead.

1742/1742 [=====] - 44s 24ms/step - loss: 0.3801 - accuracy: 0.8357 - val\_loss: 0.3581 - val\_accuracy: 0.8510

Epoch 2/10

1742/1742 [=====] - 42s 24ms/step - loss: 0.3356 - accuracy: 0.8564 - val\_loss: 0.3235 - val\_accuracy: 0.8660

Epoch 3/10

1742/1742 [=====] - 42s 24ms/step - loss: 0.3183 - accuracy: 0.8637 - val\_loss: 0.3143 - val\_accuracy: 0.8699

Epoch 4/10

1742/1742 [=====] - 43s 24ms/step - loss: 0.3043 - accuracy: 0.8715 - val\_loss: 0.3290 - val\_accuracy: 0.8629

Epoch 5/10

1742/1742 [=====] - 56s 32ms/step - loss: 0.2924 - accuracy: 0.8754 - val\_loss: 0.3011 - val\_accuracy: 0.8750

Epoch 6/10

1742/1742 [=====] - 59s 34ms/step - loss: 0.2791 - accuracy: 0.8827 - val\_loss: 0.3228 - val\_accuracy: 0.8652

Epoch 7/10

1742/1742 [=====] - 67s 38ms/step - loss: 0.2575 - accuracy: 0.8926 - val\_loss: 0.3106 - val\_accuracy: 0.8715

Epoch 8/10

1742/1742 [=====] - 62s 35ms/step - loss: 0.2280 - accuracy: 0.9065 - val\_loss: 0.3540 - val\_accuracy: 0.8653

Epoch 9/10

1742/1742 [=====] - 57s 33ms/step - loss: 0.1858 - accuracy: 0.9251 - val\_loss: 0.3895 - val\_accuracy: 0.8510


Epoch 10/10

1742/1742 [=====] - 54s 31ms/step - loss: 0.1425 - accuracy: 0.9451 - val\_loss: 0.4812 - val\_accuracy: 0.8595


Plot loss and accuracy over the training run.



```
In [10]: ▶ plt.plot(history.history['loss'])  
plt.plot(history.history['val_loss'])  
plt.legend(['Training Loss', 'Validation Loss'])  
plt.xlabel('Epoch')  
plt.ylabel('Loss')  
plt.show()
```



```
In [11]: ▶ plt.plot(history.history['accuracy'])  
plt.plot(history.history['val_accuracy'])  
plt.legend(['Training Accuracy', 'Validation Accuracy'])  
plt.xlabel('Epoch')  
plt.ylabel('Accuracy')  
plt.show()
```



Compute accuracy of the model on the training and testing sets.

```
In [12]: ► # Evaluate the model on the test data
loss, accuracy = model.evaluate(x_test, y_test)

print(f"Test Loss: {loss:.4f}")
print(f"Test Accuracy: {accuracy * 100:.2f}%")
```

```
431/431 [=====] - 4s 9ms/step - loss: 0.4732 - accuracy: 0.8609
Test Loss: 0.4732
Test Accuracy: 86.09%
```

Try a different setting to see if you can improve the test set accuracy at all. Write about the results here.

fiResults:

- I played with the VGG blueprint a bit.
- Increasing the conv sizes and the amount of conv runs actually did not improve the outcome at all. It actually showed the phenomena of overfitting discussed in class more. Because the model got too complex it fit irrelevant data
-